

Security benchmark review

103 frozen inputs · 62 public + 8 authored attacks · 33 benign · detector classification

WITHOUT LAYA / LLAMA REVIEW

85.7%

ATTACK
RECALL

60/70 attacks caught

1/33 false positives / 98.4% precision

WITH LAYA / OFFLINE OR

85.7%

ATTACK
RECALL

60/70 attacks caught

2/33 false positives / 96.8% precision

ACTUAL LAYA TRAINING

4 epochs / 192 steps · selected epoch 3

26.2M head parameters · frozen encoder

Synthetic held-out: 8/40 → 13/40 caught

0/48 benign false positives for both

REGRESSION: +0 catches / +1 false positive

SAME INPUTS / SIX DETECTOR PROFILES

PROFILE	CAUGHT	BENIGN BLOCKED
SDK regex only	0/70	2/33
Laya v0.3 only	1/70	1/33
Llama semantic review	60/70	1/33
Llama + Laya / OR	60/70	2/33
Regex + Llama / OR	60/70	3/33
Regex + Llama + Laya / OR	60/70	4/33

PUBLIC DATA / LOCAL RUNTIME

InjecAgent subset: 52/62 caught (83.9%)

Full regex base: 17/1,054 new-org; 0 SDK

Enhanced: 1,054/1,054; override prefix

Llama 3.1 8B Q4_K_M / i5-1235U CPU

LLM p50 15.44s / p95 23.47s / 0 errors

Combined profiles = offline OR. Task-aware SDK review bypasses raw regex. Laya v0.3 was not promoted for blocking.

Previously inspected local regression · Not official InjecAgent ASR or a leaderboard rank · No agent attack actions executed